

153_project

2022-04-05

```
library(tidyverse)
```

```
## -- Attaching packages ----- tidyverse 1.3.1 --
```

```
## v ggplot2 3.3.5    v purrr  0.3.4
## v tibble  3.1.6    v dplyr  1.0.8
## v tidyr   1.2.0    v stringr 1.4.0
## v readr   2.1.2    v forcats 0.5.1
```

```
## -- Conflicts ----- tidyverse_conflicts() --
```

```
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
```

```
library(TSA)
```

```
##
```

```
## Attaching package: 'TSA'
```

```
## The following object is masked from 'package:readr':
```

```
##
```

```
##      spec
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##      acf, arima
```

```
## The following object is masked from 'package:utils':
```

```
##
```

```
##      tar
```

```
library(forecast)
```

```
## Registered S3 method overwritten by 'quantmod':
```

```
##   method      from
```

```
## as.zoo.data.frame zoo
```

```
## Registered S3 methods overwritten by 'forecast':
```

```
##   method      from
```

```
## fitted.Arima TSA
```

```
## plot.Arima   TSA
```

```
library(astsa)
```


```
##  
## Attaching package: 'astsa'  
  
## The following object is masked from 'package:forecast':  
##  
##     gas
```

```
source('cleaning.R')
```


```
##  
## Attaching package: 'lubridate'  
  
## The following objects are masked from 'package:base':  
##  
##     date, intersect, setdiff, union  
  
## New names:  
## * ' ' -> ...1  
  
## Rows: 1746661 Columns: 29  
## -- Column specification -----  
## Delimiter: ","  
## chr   (8): Address, State, County, City, NO2 Units, O3 Units, SO2 Units, CO ...  
## dbl  (20): ...1, State Code, County Code, Site Num, NO2 Mean, NO2 1st Max Va...  
## date  (1): Date Local  
##  
## i Use 'spec()' to retrieve the full column specification for this data.  
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.  
## 'summarise()' has grouped output by 'year'. You can override using the '.groups' argument.
```

Original data

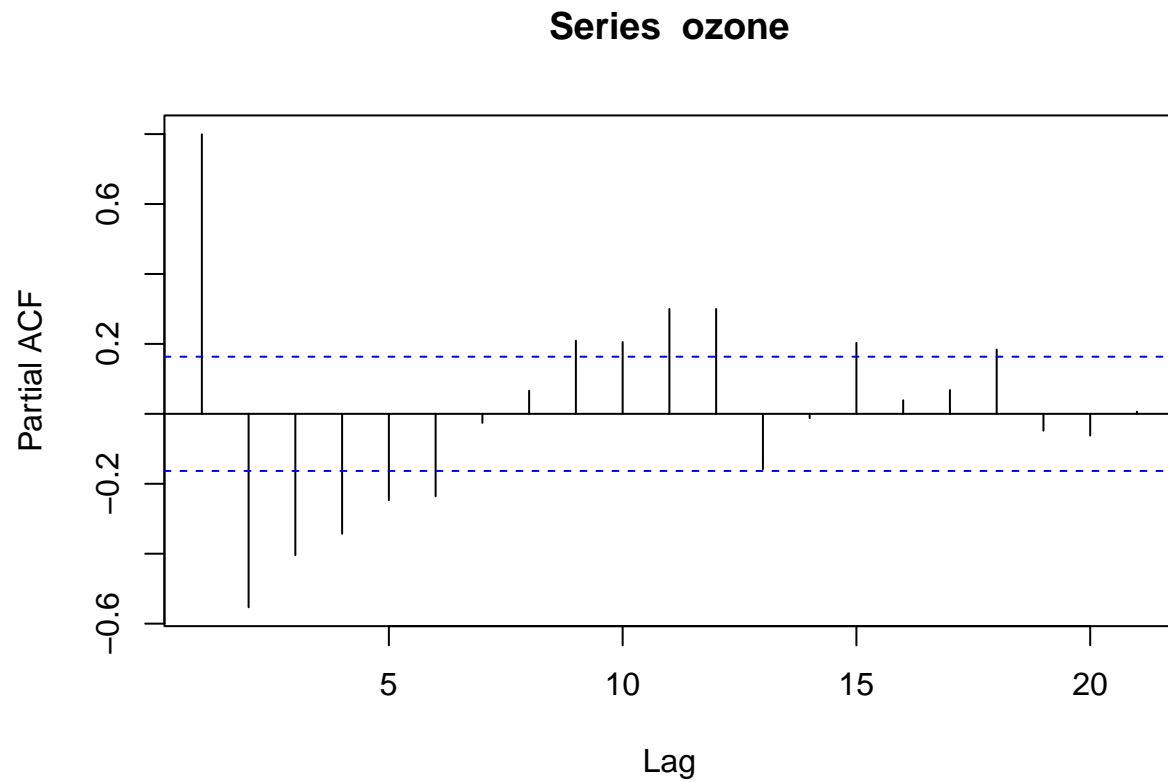
```
ozone <- phoenix$o3  
ozone = ozone - mean(ozone) # mean centered  
plot(ozone, type = "l")
```



```
acf(ozone, lag.max = 30)
```




```
pacf(ozone)
```



1. The original data have clear sign of seasonality, but there seems to be linear upward trend. 2. The ACF PACF plot shows the strong sign of seasonality

Sinusoidal fitting

```
t = 1:length(ozone)
# Check the periodogram
periodo = periodogram(ozone,plot=TRUE,ylab="Periodogram", xlab="Frequency") # There is one significant
```




```
# Get the high magnitudes in descending order
order_spec = sort(periodo$spec,decreasing = TRUE)

# Get the frequency that gives max magnitude
first_max = order_spec[1]
first_maximizing_freq = periodo$freq[periodo$spec==first_max]
first_sin_max = sin(2*pi*first_maximizing_freq*t)
first_cos_max = cos(2*pi*first_maximizing_freq*t)


# Max Sinusoidal fitting
ozone_sinusoid_model = lm(ozone ~ first_sin_max*(1+t)+first_cos_max*(1+t))
print(ozone_sinusoid_model$coefficients)

##      (Intercept)  first_sin_max          t  first_cos_max first_sin_max:t
## -3.009489e-03  -1.160248e-03   4.163369e-05  -1.233604e-02   1.963185e-06
## t:first_cos_max
## -1.058031e-05

# Overlay the sinusoidal fitting over the original plot
plot(ozone,type = "l")
lines(t,ozone_sinusoid_model$fitted.values,col = "red")
```




```
# Get the residual, hoping for removing seasonality  
ozone_sinusoid_residual = ozone_sinusoid_model$residuals  
plot(ozone_sinusoid_residual, type = "l") # residual seems to be stationary
```




```
stats::acf(ozone_sinusoid_residual)
```


Series ozone_sinusoid_residual



```
stats::pacf(ozone_sinusoid_residual)
```

Series ozone_sinusoid_residual



$$f(t) = -0.003 + 0.000042 * t - 0.0012 * \sin(t) - 0.0123 * \cos(t) + 0.000002 * t * \sin(t) - 0.000011 * t * \cos(t)$$

1. There is one significant peak in the periodogram
2. The residual plot shows some seasonality but the plot seems to be AR process and possible seasonal ARMA

Residual Fitting : SARIMA(2,0,0)(1,0,1)12


```
model1 <- sarima(ozone_sinusoid_residual, p=2, d=0, q=0, P=1, D=0, Q=1, S=12) # fit the model
```

```
## initial value -5.921137
## iter 2 value -6.020457
## iter 3 value -6.079641
## iter 4 value -6.085871
## iter 5 value -6.091190
## iter 6 value -6.095582
## iter 7 value -6.095958
## iter 8 value -6.095969
## iter 9 value -6.095970
## iter 9 value -6.095970
## iter 9 value -6.095970
## final value -6.095970
## converged
```

```

## initial value -6.075290
## iter 2 value -6.078754
## iter 3 value -6.081474
## iter 4 value -6.085882
## iter 5 value -6.088012
## iter 6 value -6.092257
## iter 7 value -6.095296
## iter 8 value -6.097651
## iter 9 value -6.099009
## iter 10 value -6.100678
## iter 11 value -6.101042
## iter 12 value -6.101155
## iter 13 value -6.101161
## iter 14 value -6.101163
## iter 15 value -6.101163
## iter 16 value -6.101166
## iter 17 value -6.101167
## iter 18 value -6.101167
## iter 19 value -6.101168
## iter 20 value -6.101168
## iter 21 value -6.101169
## iter 22 value -6.101169
## iter 22 value -6.101169
## iter 22 value -6.101169
## final value -6.101169
## converged

```



```
coeff_table <- as.data.frame(model1$tttable)
coeff_table <-coeff_table %>% mutate(ci_lower = Estimate-1.96*SE,ci_upper = Estimate+1.96*SE)
coeff_table # show estimated coefficient and its ci
```


```
##      Estimate      SE t.value p.value ci_lower ci_upper
## ar1      0.0903 0.0837  1.0792  0.2824 -0.073752  0.254352
## ar2     -0.0160 0.0881 -0.1810  0.8566 -0.188676  0.156676
## sar1      0.9462 0.0545 17.3486  0.0000  0.839380  1.053020
## sma1     -0.7176 0.1439 -4.9858  0.0000 -0.999644 -0.435556
## xmean    -0.0002 0.0006 -0.3728  0.7099 -0.001376  0.000976
```

Residual Fitting : SARIMA(1,0,0)(1,0,1)12

```
model2 <- sarima(ozone_sinusoid_residual, p=1, d=0, q=0, P=1, D=0, Q=1, S=12) # fit the model
```

```
## initial value -5.924525
## iter 2 value -6.014280
## iter 3 value -6.072474
## iter 4 value -6.080884
## iter 5 value -6.087853
## iter 6 value -6.094164
## iter 7 value -6.094618
## iter 8 value -6.094677
## iter 9 value -6.094683
## iter 10 value -6.094683
## iter 11 value -6.094684
## iter 12 value -6.094684
## iter 12 value -6.094684
## iter 12 value -6.094684
## final value -6.094684
## converged
## initial value -6.075303
## iter 2 value -6.078312
## iter 3 value -6.083138
## iter 4 value -6.089196
## iter 5 value -6.092910
## iter 6 value -6.095028
## iter 7 value -6.096554
## iter 8 value -6.098751
## iter 9 value -6.100853
## iter 10 value -6.101030
## iter 11 value -6.101038
## iter 12 value -6.101038
## iter 13 value -6.101041
## iter 14 value -6.101048
## iter 15 value -6.101050
## iter 16 value -6.101052
## iter 17 value -6.101053
## iter 18 value -6.101054
## iter 19 value -6.101055
```

```
## iter 20 value -6.101055
## iter 20 value -6.101055
## iter 20 value -6.101055
## final value -6.101055
## converged
```



```
coeff_table <- as.data.frame(model2$table)
coeff_table <- coeff_table %>% mutate(ci_lower = Estimate-1.96*SE, ci_upper = Estimate+1.96*SE)
coeff_table # show estimated coefficient and its ci
```

	Estimate	SE	t.value	p.value	ci_lower	ci_upper
## ar1	0.0888	0.0832	1.0667	0.2879	-0.074272	0.251872
## sar1	0.9489	0.0498	19.0430	0.0000	0.851292	1.046508
## sma1	-0.7226	0.1379	-5.2393	0.0000	-0.992884	-0.452316
## xmean	-0.0002	0.0006	-0.3640	0.7164	-0.001376	0.000976

Residual Fitting : SARIMA(2,1,1)(1,0,1)₁₂


```
model3 <- sarima(ozone_sinusoid_residual, p=2, d=1, q=1, P=1, D=0, Q=1, S=12) # fit the model
```

```
## initial value -5.626415
## iter 2 value -5.817904
```

```

## iter    3 value -5.967005
## iter    4 value -6.000235
## iter    5 value -6.022716
## iter    6 value -6.038025
## iter    7 value -6.061036
## iter    8 value -6.062258
## iter    9 value -6.063975
## iter   10 value -6.065519
## iter   11 value -6.066125
## iter   12 value -6.066443
## iter   13 value -6.067328
## iter   14 value -6.067686
## iter   15 value -6.067780
## iter   16 value -6.067831
## iter   17 value -6.067832
## iter   18 value -6.067883
## iter   19 value -6.067920
## iter   20 value -6.067946
## iter   21 value -6.067967
## iter   22 value -6.068029
## iter   23 value -6.068059
## iter   24 value -6.068062
## iter   25 value -6.068063
## iter   25 value -6.068063
## iter   25 value -6.068063
## final   value -6.068063
## converged
## initial  value -6.049027
## iter    2 value -6.063938
## iter    3 value -6.069719
## iter    4 value -6.077660
## iter    5 value -6.081289
## iter    6 value -6.083983
## iter    7 value -6.086264
## iter    8 value -6.089727
## iter    9 value -6.093924
## iter   10 value -6.097960
## iter   11 value -6.098795
## iter   12 value -6.099032
## iter   13 value -6.099192
## iter   14 value -6.099209
## iter   15 value -6.099215
## iter   16 value -6.099216
## iter   17 value -6.099217
## iter   18 value -6.099217
## iter   19 value -6.099217
## iter   20 value -6.099218
## iter   21 value -6.099219
## iter   22 value -6.099220
## iter   23 value -6.099220
## iter   23 value -6.099220
## iter   23 value -6.099220
## final   value -6.099220
## converged

```



```
coeff_table <- as.data.frame(model3$tttable)
coeff_table <-coeff_table %>% mutate(ci_lower = Estimate-1.96*SE,ci_upper = Estimate+1.96*SE)
coeff_table # show estimated coefficient and its ci
```

	Estimate	SE	t.value	p.value	ci_lower	ci_upper
## ar1	0.0192	0.0892	0.2150	0.8301	-0.155632	0.194032
## ar2	-0.0791	0.0894	-0.8846	0.3779	-0.254324	0.096124
## ma1	-0.9352	0.0317	-29.5144	0.0000	-0.997332	-0.873068
## sar1	0.9602	0.0366	26.2225	0.0000	0.888464	1.031936
## sma1	-0.7544	0.1106	-6.8202	0.0000	-0.971176	-0.537624
## constant	0.0000	0.0001	-0.0145	0.9884	-0.000196	0.000196

Residual Fitting : SARIMA(2,0,1)(1,0,1)₁₂

```
model4 <- sarima(ozone_sinusoid_residual, p=2, d=0, q=1, P=1, D=0, Q=1, S=12) # fit the model
```

```
## initial value -5.921137
## iter 2 value -6.021739
## iter 3 value -6.079075
## iter 4 value -6.085438
## iter 5 value -6.091319
## iter 6 value -6.095255
```


```

## iter    7 value -6.095886
## iter    8 value -6.095982
## iter    9 value -6.095984
## iter   10 value -6.095987
## iter   11 value -6.095993
## iter   12 value -6.096006
## iter   13 value -6.096024
## iter   14 value -6.096036
## iter   15 value -6.096039
## iter   16 value -6.096040
## iter   17 value -6.096040
## iter   18 value -6.096040
## iter   19 value -6.096040
## iter   19 value -6.096040
## iter   19 value -6.096040
## final   value -6.096040
## converged
## initial  value -6.075372
## iter    2 value -6.077112
## iter    3 value -6.080624
## iter    4 value -6.083981
## iter    5 value -6.086061
## iter    6 value -6.091109
## iter    7 value -6.094569
## iter    8 value -6.096817
## iter    9 value -6.097887
## iter   10 value -6.099778
## iter   11 value -6.100684
## iter   12 value -6.101115
## iter   13 value -6.101118
## iter   14 value -6.101118
## iter   15 value -6.101118
## iter   16 value -6.101124
## iter   17 value -6.101136
## iter   18 value -6.101149
## iter   19 value -6.101171
## iter   20 value -6.101187
## iter   21 value -6.101190
## iter   22 value -6.101192
## iter   23 value -6.101197
## iter   24 value -6.101201
## iter   25 value -6.101209
## iter   26 value -6.101219
## iter   27 value -6.101230
## iter   28 value -6.101233
## iter   29 value -6.101237
## iter   29 value -6.101237
## iter   29 value -6.101237
## final   value -6.101237
## converged

## Warning in sqrt(diag(fitit$var.coef)): NaNs produced

## Warning in sqrt(diag(fitit$var.coef)): NaNs produced

```

```
coeff_table <- as.data.frame(model4$tttable)
coeff_table <-coeff_table %>% mutate(ci_lower = Estimate-1.96*SE,ci_upper = Estimate+1.96*SE)
coeff_table # show estimated coefficient and its ci
```

	Estimate	SE	t.value	p.value	ci_lower	ci_upper
## ar1	-0.3458	NaN	NaN	NaN	NaN	NaN
## ar2	0.0214	NaN	NaN	NaN	NaN	NaN
## ma1	0.4374	NaN	NaN	NaN	NaN	NaN
## sar1	0.9498	0.0488	19.4654	0.0000	0.854152	1.045448
## sma1	-0.7278	0.1344	-5.4138	0.0000	-0.991224	-0.464376
## xmean	-0.0002	0.0006	-0.3604	0.7191	-0.001376	0.000976

Evaluation Matrix

```
# AIC, AICc, BIC
eval<- function(model){
  return (c(model$AIC, model$AICc,model$BIC))
}

m1_evaludation = eval(model1)
m2_evaludation = eval(model2)
m3_evaludation = eval(model3)
m4_evaludation = eval(model4)
```

```
eval_matrix = rbind(m1_evaluation,m2_evaluation,m3_evaluation,m4_evaluation)
rownames(eval_matrix) = c("SARIMA(2,0,0)(1,0,1)12", "SARIMA(1,0,0)(1,0,1)12",
                          "SARIMA(2,1,1)(1,0,1)12 ", "SARIMA(2,0,1)(1,0,1)12 ")
colnames(eval_matrix) = c("AIC", "AICc", "BIC")
eval_matrix
```

```
##              AIC      AICc      BIC
## SARIMA(2,0,0)(1,0,1)12 -9.281128 -9.278109 -9.157386
## SARIMA(1,0,0)(1,0,1)12 -9.294789 -9.292790 -9.191670
## SARIMA(2,1,1)(1,0,1)12 -9.262661 -9.258342 -9.117626
## SARIMA(2,0,1)(1,0,1)12 -9.267376 -9.263118 -9.123010
```

The first two models, SARIMA(2,0,0)(1,0,1)12 , SARIMA(1,0,0)(1,0,1)12 gives the lowest AIC,AICc,BIC

Cross Validation


Train : 2004 ~ 2011 Test : 2012 - 2015


```
sse1 = c()
sse2 = c()
test_years = seq(12,15,1)
for (year in test_years) {


train_index = 1:(12*(year-4))
test_index = (12*(year-4)+1):(12*(year-4+1))


train <- ozone_sinusoid_residual[train_index]
test <- ozone_sinusoid_residual[test_index]
m1_forecast <- sarima.for(train, p=2, d=0, q=0, P=1, D=0, Q=1, S=12, n.ahead=12)$pred
m2_forecast <- sarima.for(train, p=1, d=0, q=0, P=1, D=0, Q=1, S=12, n.ahead=12)$pred


sse1 = c(sse1,sum((m1_forecast - test)^2))
sse2 = c(sse2,sum((m2_forecast - test)^2)) }
```




















```
sse = rbind(sum(sse1),sum(sse2))
rownames(sse) = c("SARIMA(2,0,0)(1,0,1)12","SARIMA(1,0,0)(1,0,1)12")
colnames(sse) = c("SSE")
print(sse) # The SSE for both model are low, but SARIMA(2,0,0)(1,0,1)12 is slightly better
```

```
##                               SSE
## SARIMA(2,0,0)(1,0,1)12 0.0002583248
## SARIMA(1,0,0)(1,0,1)12 0.0002598963
```